



TLS 1.3 Handshake In Linux Kernel

Xin Long
Aug, 2022

Overview

APIs & Usage:

```
TLS_HS (msg)
|-- TLS_HS_GEN (sk)
|  |
|  |-- TCP (TLS_HS_TCP)
|  |-- NFS
|
|-- QUIC
```

Implementation:

```
repo: https://github.com/lxin/tls\_hs
- include/crypto/tls_hs.h
- crypto/tls_hs.c
(CONFIG_CRYPTOTLS_HS=m/y)
```

APIs: TLS_HS Core (msg)

`tls = tls_handshake_create(is_serv)`

`tls_handshake_get/set(tls, opt, vec)`

- opt:

- ▶ TLS_T_PKEY: private key.
- ▶ TLS_T_PSK:
- ▶ TLS_T_CA:
- ▶ TLS_T_CRT_REQ: client certificate request.
- ▶ TLS_T_CRT: certificate chain.
- ▶ TLS_T_EXT: unknown extension.
- ▶ TLS_T_EARLY: flag for early data send/recv.

`state = tls_handshake(tls, msg)`

- msg:

- ▶ **input:** a tls 1.3 msg to process.
- ▶ **output:** a tls msg to send for reply.

Note: when input is NULL, output is tls 1.3 CH/EE msg.

- state:

- ▶ TLS_ST_START
- ▶ TLS_ST_RCVD
- ▶ TLS_ST_WAIT
- ▶ TLS_ST_CONNECTED
- ▶ < 0: Error

APIs: TLS_HS Extra (msg)

- `tls_secret_get(tls, level, srt)`
- `tls_hkdf_expand(tls, srt, h, l, k)`
- `tls_hkdf_extract(tls, srt, h, k)`

- level:

- ▶ `TLS_SE_RMS`: resumption secret.
- ▶ `TLS_SE_EA`: early secret.
- ▶ `TLS_SE_HS`: handshake secret.
- ▶ `TLS_SE_AP`: master secret.

- `tls_handshake_destroy(tls)`
- `tls_handshake_post(tls, type, msg)`

- type:

- ▶ `TLS_P_NONE`: process post the handshake msg.
- ▶ `TLS_P_TICKET`: create session ticket msg.
- ▶ `TLS_P_KEY_UPDATE`: create new keys and key update msg.

APIs: TLS_HS_GEN Core (sk)

tls = `tls_sk_handshake`(sk, data, keyring, flag)

- sk:

- ▶ **input:** a TCP established socket.
- ▶ **output:** a kTLS socket with keys set.

- data:

- ▶ early data to send and early data received.

- keyring:

- ▶ set for reading keys/crts from userspace, and left none if the keys/crts are set by kernel users via `tls_handshake_set(...)`.

- flag:

- ▶ `TLS_F_SERV`: works as a server.
- ▶ `TLS_F_PSK/CRT/CRT_REQ`:
- ▶ `TLS_F_NO_KTLS`:
w/o kTLS users can use `tls_ap_de/encrypt()` to send/recv app data.

- **tls (return obj):**

- ▶ users can either save it for future post handshake msg processing or destroy it.

APIs: TLS_HS_GEN Extra (sk)

- `tls_sk_handshake_post(sk, tls, type, msg)`
- `tls_ap_encrypt(tls, data, seq)`
- `tls_ap_decrypt(tls, data, seq)`

Usage: QUIC

Module in Kernel:

- ▶ repo: https://github.com/lxin/tls_hs/tree/quic
 - net/quic/*
 - include/net/quic/quic.h
 - include/uapi/linux/quic.h
- (CONFIG_IP_QUIC=m/y)

Kernel Code:

- ▶ handshake code: https://github.com/lxin/tls_hs/blob/quic/net/quic/frame.c#L654

User Programs:

- ▶ https://github.com/lxin/tls_hs-tests#ii-quic

Usage: TCP/NFS

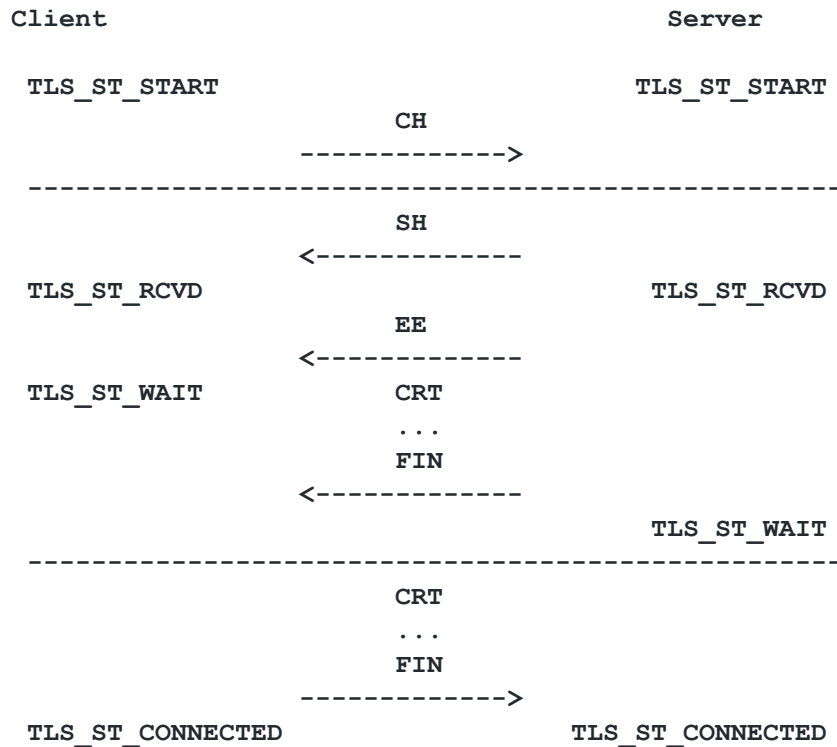
Kernel Code:

- ▶ TCP: https://github.com/ixn/tls_hs/blob/tcp/net/ipv4/tcp.c#L4363
- ▶ NFS: https://github.com/ixn/tls_hs/blob/sunrpc/net/sunrpc/xprtsock.c#L2433

User Programs:

- ▶ https://github.com/ixn/tls_hs-tests#i-tls_hs_tcp

Implementation: state transition



Implementation: ciphers & algorithms

- ▶ HKDF extract/expand: hmac(sha256)
- ▶ ECDH key exchange: secp256r1(0x0017)
- ▶ PSK exchange mode: psk_dhe_ke(1)
- ▶ Certificate: rsa_pkcs1_sha256(0x0401)/rsa_pss_rsae_sha256(0x0804)
- ▶ Signature Algorithm: rsa_pss_rsae_sha256(0x0804)
- ▶ AEAD: TLS_AES_128_GCM_SHA256(0x1301)

Implementation: functions

- ▶ Certificate Chain and CA on Both Sides
- ▶ PSK
- ▶ Session Resumption
- ▶ Early Data
- ▶ Hello Retry Request on Server

Implementation: messages & extensions

-	TLS_MT_HELLO_RETRY_REQUEST	0
-	TLS_MT_CLIENT_HELLO	1
-	TLS_MT_SERVER_HELLO	2
-	TLS_MT_NEWSESSION_TICKET	4
-	TLS_MT_END_OF_EARLY_DATA	5
-	TLS_MT_ENCRYPTED_EXTENSIONS	8
-	TLS_MT_CERTIFICATE	11
-	TLS_MT_CERTIFICATE_REQUEST	13
-	TLS_MT_CERTIFICATE_VERIFY	15
-	TLS_MT_FINISHED	20
-	TLS_MT_KEY_UPDATE	24

-	TLS_EXT_server_name	0
-	TLS_EXT_supported_groups	10
-	TLS_EXT_ec_point_formats	11
-	TLS_EXT_signature_algorithms	13
-	TLS_EXT_heartbeat	15
-	TLS_EXT_alpn	16
-	TLS_EXT_signed_cert_timestamp	18
-	TLS_EXT_padding	21
-	TLS_EXT_encrypt_then_mac	22
-	TLS_EXT_extended_master_srt	23
-	TLS_EXT_session_ticket	35
-	TLS_EXT_psk	41
-	TLS_EXT_early_data	42
-	TLS_EXT_supported_versions	43
-	TLS_EXT_cookie	44
-	TLS_EXT_psk_kex_modes	45
-	TLS_EXT_certificate_authorities	47
-	TLS_EXT_post_handshake_auth	49
-	TLS_EXT_signature_algs_cert	50
-	TLS_EXT_key_share	51

Implementation: memory management

struct tls_hs

Creating and Parsing msgs:

- ▶ struct tls_hello h;
- ▶ struct tls_vec ext; /* 1 page */
- ▶ struct tls_vec cmsg; /* 1 page */
- ▶ struct tls_vec omsg; /* 1 page */

Secrets and HS msg buffers:

- ▶ struct tls_vec buf[TLS_B_MAX];
- ▶ struct tls_vec srt[TLS_SE_MAX];

State and Flags:

- ▶ u8 state:2, early:1, is_serv:1, crt_req:1;

Crypto API algorithm objects:

- ▶ struct crypto_kpp *kpp_tfm;
- ▶ struct crypto_aead *aead_tfm;
- ▶ struct crypto_shash *srt_tfm;
- ▶ struct crypto_shash *hash_tfm;
- ▶ struct crypto_akcipher *akc_tfm;

Certs and Keys:

- ▶ struct tls_vec pkey;
- ▶ struct tls_crt *tcrt;
- ▶ struct tls_crt *rcrt;
- ▶ struct tls_crt *ca;
- ▶ struct tls_psk *psks;

END